
🏥 Capstone Project: Healthcare Data Engineering Platform on Azure Databricks

---

🎯 Objective

Build a complete data engineering solution for a fictional healthcare analytics company — **MediPulse Analytics** — using **Azure Databricks**. You will ingest, clean, transform, and analyze healthcare data from multiple sources, store it using Delta Lake, implement incremental loads, and prepare analytical views for downstream machine learning and dashboards.

📌 Project Scenario

Company: MediPulse Analytics **Goal:** To process and analyze healthcare patient data, hospital data, and appointment records for building KPIs and dashboards.

📂 Data Sources

You will simulate three raw datasets:

1️⃣ `patients.csv`

```
patient_id,name,age,gender,region
P001,Arjun Mehta,34,M,North
P002,Neha Sharma,29,F,South
P003,Rahul Gupta,40,M,East
P004,Sneha Nair,25,F,West
```

2️⃣ `hospitals.json`

```
[
  {"hospital_id": "H001", "hospital_name": "City Care", "region": "North"},
  {"hospital_id": "H002", "hospital_name": "LifePlus", "region": "South"},
  {"hospital_id": "H003", "hospital_name": "MediHope", "region": "East"},
  {"hospital_id": "H004", "hospital_name": "CureWell", "region": "West"}
]
```

3️⃣ `appointments_day1.csv`

```
appointment_id,patient_id,hospital_id,appointment_date,diagnosis,cost,status
A1001,P001,H001,2024-01-10,Diabetes,400,Completed
A1002,P002,H002,2024-01-11,Flu,250,Completed
A1003,P003,H003,2024-01-11,Heart Disease,1000,Pending
A1004,P004,H004,2024-01-12,Allergy,300,Completed
```

🚀 Step 1 — Bronze Layer: Raw Ingestion

 Tasks:

- Read CSV and JSON data into DataFrames.
- Write them as Delta tables (`bronze_patients` , `bronze_hospitals` , `bronze_appointments`).

Step 2 — Silver Layer: Data Cleansing & Transformation

 Tasks:

- Filter out `Pending` appointments.
- Join patients and hospitals to enrich appointment data.
- Add new calculated column: `year = year(appointment_date)` and `month` .
- Store output as `silver_appointments` .

Step 3 — Gold Layer: Analytical Aggregations

 Tasks:

- Total revenue per hospital.
- Total patients per region.
- Top 3 most expensive diagnosis categories.
- Store as `gold_healthcare_summary` .

Step 4 — Incremental Load Simulation

 Tasks:

- Create `appointments_day2.csv` with new data.
- Use `MERGE` or `Upsert` to update the silver table.
- Show how incremental data changes the gold table.

Step 5 — Delta Lake Features

 Tasks:

- Use **Time Travel** to view the gold table before incremental load.
- Use **Vacuum** to clean up historical versions.
- Use **Optimize + Z-Ordering** on `hospital_id` .

Analytical Questions to Solve

1. Total revenue generated by each hospital.
2. Average cost per diagnosis category.
3. Number of patients served per region.
4. Trend of appointments month-over-month.
5. Top 5 most expensive treatments in the last 6 months.
